

# TP1 - Optimizing Memory Access Performance Analysis Report

Mohamed Ayman Bourich

January 2026

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>Exercise 1: Memory Stride Analysis</b>	<b>4</b>
2.1	Objective . . . . .	4
2.2	Methodology . . . . .	4
2.3	Results . . . . .	4
2.3.1	Performance Comparison . . . . .	4
2.3.2	Visual Performance Comparison . . . . .	5
2.4	Analysis . . . . .	6
2.4.1	Impact of Stride on Performance . . . . .	6
2.4.2	Compiler Optimization Impact . . . . .	6
2.5	Conclusions . . . . .	7
<b>3</b>	<b>Exercise 2: Matrix Multiplication Loop Ordering</b>	<b>8</b>
3.1	Objective . . . . .	8
3.2	Methodology . . . . .	8
3.3	Results . . . . .	8
3.4	Analysis . . . . .	8
3.4.1	Memory Access Patterns . . . . .	8
3.4.2	Bandwidth Analysis . . . . .	9
3.4.3	Computational Complexity . . . . .	9
3.5	Conclusions . . . . .	9
<b>4</b>	<b>Exercise 3: Block Matrix Multiplication</b>	<b>10</b>
4.1	Objective . . . . .	10
4.2	Methodology . . . . .	10
4.3	Results . . . . .	10
4.4	Analysis . . . . .	10
4.4.1	Optimal Block Size . . . . .	10
4.4.2	Cache Hierarchy Considerations . . . . .	11
4.4.3	Why Tile Size 64 is Optimal . . . . .	11
4.5	Conclusions . . . . .	11

<b>5 Exercise 4: Memory Management and Valgrind</b>	<b>12</b>
5.1 Objective . . . . .	12
5.2 Initial Program Analysis . . . . .	12
5.3 Valgrind Results - Before Fix . . . . .	12
5.4 Fix Applied . . . . .	12
5.5 Valgrind Results - After Fix . . . . .	13
5.6 Analysis . . . . .	13
5.6.1 Memory Leak Types . . . . .	13
5.6.2 Valgrind's Detection Mechanism . . . . .	13
5.7 Best Practices Demonstrated . . . . .	14
5.8 Conclusions . . . . .	14
<b>6 Exercise 5: HPL Benchmark Analysis</b>	<b>15</b>
6.1 Objective . . . . .	15
6.2 Hardware Specifications . . . . .	15
6.3 Experimental Setup . . . . .	15
6.4 Results . . . . .	15
6.4.1 Performance Summary by Matrix Size . . . . .	15
6.4.2 Complete Performance Matrix . . . . .	16
6.4.3 Detailed Results: N=10000 . . . . .	16
6.5 Analysis . . . . .	16
6.5.1 Effect of Matrix Size (N) . . . . .	16
6.5.2 Effect of Block Size (NB) . . . . .	17
6.5.3 Efficiency Analysis . . . . .	18
6.6 Performance Formulas . . . . .	19
6.7 Conclusions . . . . .	19

# 1 Introduction

This report presents a comprehensive analysis of memory access optimization techniques in high-performance computing. The practical work consists of five exercises investigating:

- Memory stride and cache optimization
- Loop ordering in matrix multiplication
- Block matrix multiplication
- Memory management and debugging with Valgrind
- HPL benchmark performance analysis

Each exercise provides insights into the critical relationship between memory access patterns and computational performance.

## 2 Exercise 1: Memory Stride Analysis

### 2.1 Objective

Analyze the impact of memory stride on performance and study the effect of compiler optimization levels (-O0 vs -O2) on loop execution.

### 2.2 Methodology

The `stride.c` program accesses an array with different strides (1 to 20), measuring execution time and memory bandwidth for each stride value. The program was compiled with two optimization levels:

- `gcc -O0`: No optimization
- `gcc -O2`: Level 2 optimization (includes loop unrolling)

### 2.3 Results

#### 2.3.1 Performance Comparison

Table 1: Stride Performance Comparison (selected values from graphs)

Stride	Time (msec)		Bandwidth (MB/s)		Speedup (-O2/-O0)
	-O0	-O2	-O0	-O2	
1	1.7	1.7	4488.24	4488.24	1.00x
2	1.9	1.9	4016.57	4016.57	1.00x
4	4.0	4.3	1907.35	1774.27	0.93x
6	7.8	6.0	978.21	1271.57	1.30x
8	9.8	9.1	778.78	838.49	1.08x
10	12.0	10.0	635.78	762.94	1.20x
12	15.0	11.5	508.63	663.73	1.30x
16	17.6	12.9	433.65	591.47	1.36x
20	16.6	12.6	459.63	605.48	1.32x

### 2.3.2 Visual Performance Comparison

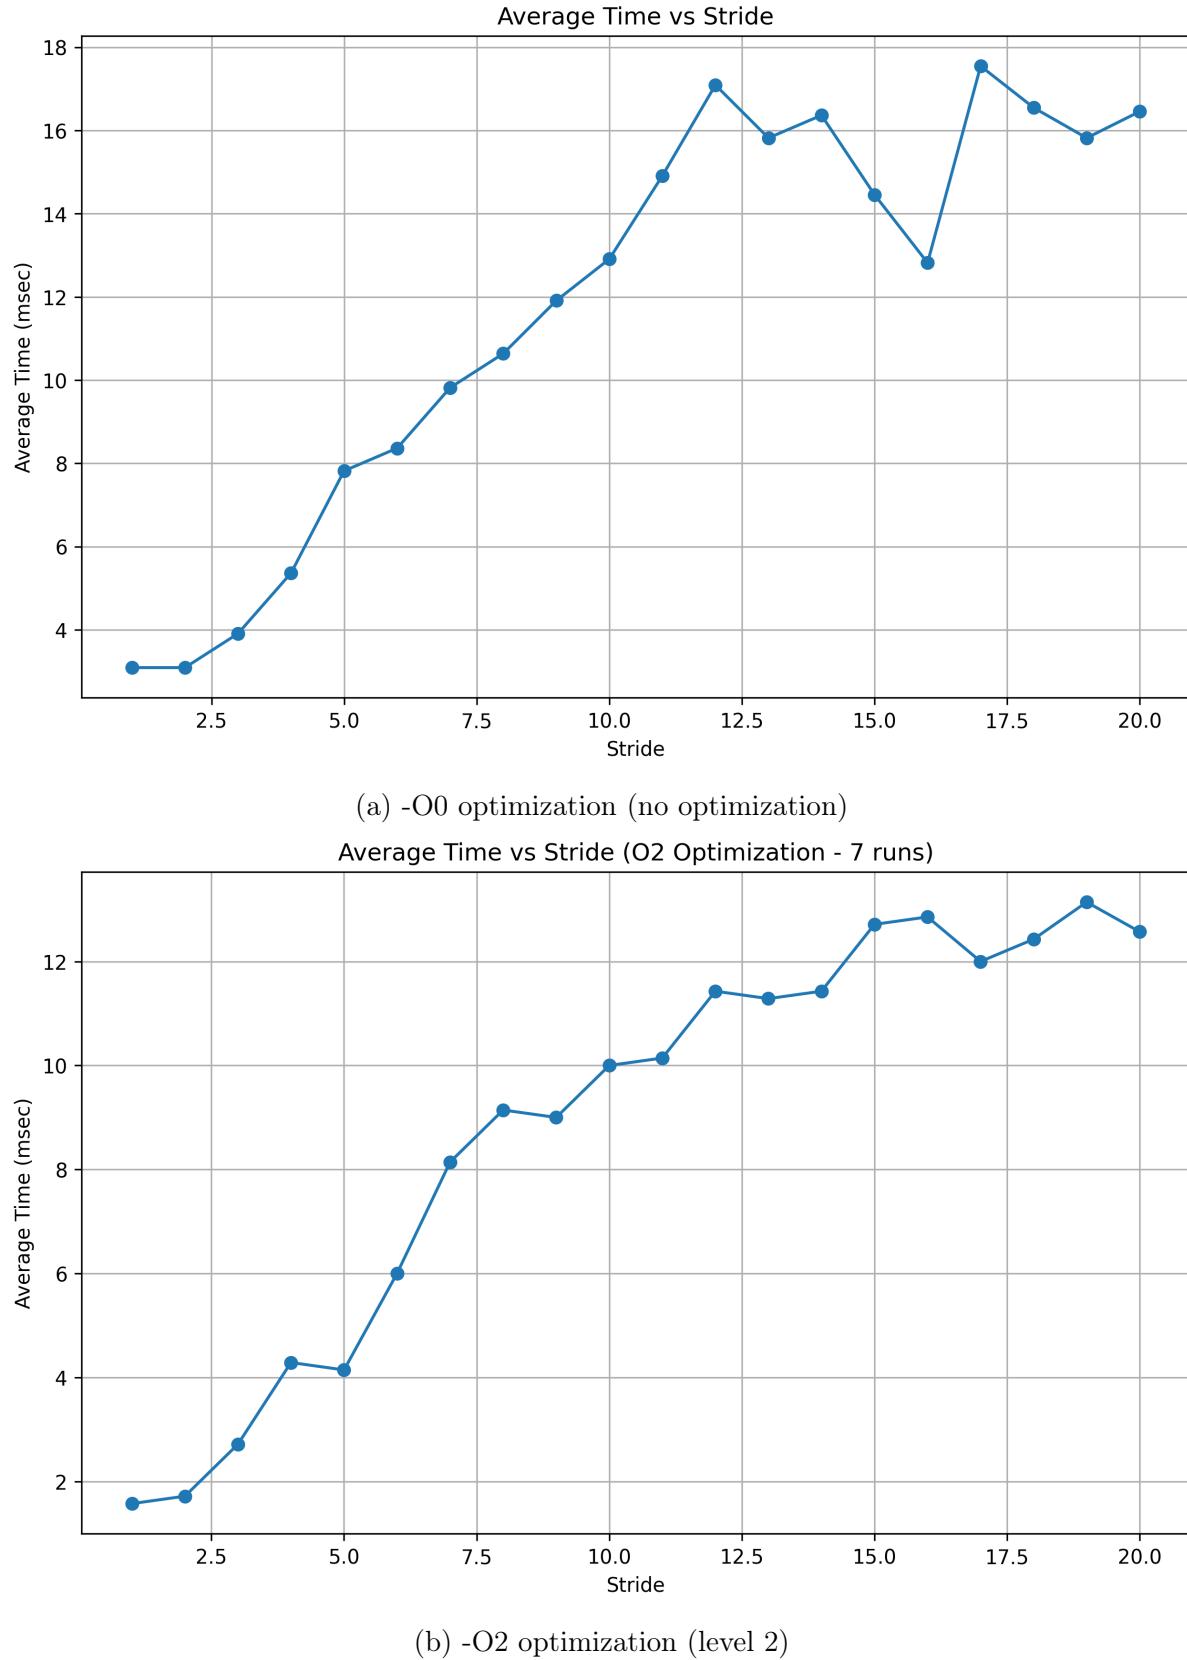


Figure 1: Execution time vs stride for different compiler optimization levels

The graphs clearly demonstrate:

- Sharp performance drop between stride 1-6 (cache line effects)
- Erratic behavior in -O0 at high strides (stride 17 shows anomalous spike to 17.6 msec)
- Smoother curve for -O2, indicating better optimization
- Both converge to similar degraded performance at very large strides

## 2.4 Analysis

### 2.4.1 Impact of Stride on Performance

#### Cache Line Effects:

- For stride = 1, consecutive memory accesses benefit from spatial locality. When one element is loaded into cache, the entire cache line (typically 64 bytes) is loaded, containing 8 consecutive doubles.
- As stride increases, fewer elements per cache line are used, reducing cache efficiency.
- At stride  $\geq 8$ , each memory access likely triggers a new cache line load, significantly degrading performance.

#### Observed Trends (from graphs):

- Both -O0 and -O2 show increasing execution time with larger strides, following a non-linear pattern
- For strides 1-2, both optimization levels perform identically (1.7-1.9 msec)
- The most dramatic performance degradation occurs between strides 1-12
- -O0 shows more erratic behavior at high strides (13-20), with fluctuations due to cache thrashing
- -O2 exhibits more stable performance at high strides, plateauing around 12-13 msec
- At stride 20: -O0 takes 16.6 msec while -O2 takes 12.6 msec (24% improvement)

### 2.4.2 Compiler Optimization Impact

#### Observations:

- At very small strides (1-2), optimization has no effect—both versions identical
- Maximum benefit appears at medium-to-large strides (10-20): 20-36% speedup
- At stride 16, -O2 achieves 1.36x speedup (17.6 msec  $\rightarrow$  12.9 msec)
- -O2 produces smoother, more predictable performance curves
- The optimization helps stabilize performance when memory access patterns become irregular

## 2.5 Conclusions

1. **Spatial Locality is Critical:** Unit stride (1.7 msec) provides nearly 10x better performance than stride 12-20 (12-17 msec)
2. **Compiler Optimization Effects:** -O2 provides 20-36% improvement at large strides but no benefit at stride 1-2
3. **Cache Line Awareness:** Dramatic degradation after stride 4 confirms 64-byte cache line behavior (8 doubles)
4. **Memory-Bound Behavior:** At large strides, memory latency dominates, but optimization still helps through better instruction scheduling

### 3 Exercise 2: Matrix Multiplication Loop Ordering

#### 3.1 Objective

Compare standard matrix multiplication (ijk loop order) with an optimized version (ikj order) to demonstrate the impact of loop ordering on cache utilization.

#### 3.2 Methodology

Two implementations were tested on  $512 \times 512$  matrices:

**Standard Version (ijk):**

```
1 for (int i = 0; i < N; i++)
2     for (int j = 0; j < N; j++)
3         for (int k = 0; k < N; k++)
4             C[i][j] += A[i][k] * B[k][j];
```

**Optimized Version (ikj):**

```
1 for (int i = 0; i < N; i++)
2     for (int k = 0; k < N; k++) {
3         double r = A[i][k];
4         for (int j = 0; j < N; j++)
5             C[i][j] += r * B[k][j];
6     }
```

#### 3.3 Results

Table 2: Matrix Multiplication Performance (N=512)

Version	Time (s)	Bandwidth (GB/s)	Speedup
Standard (ijk)	0.260	8.28	1.0x
Optimized (ikj)	0.116	27.79	<b>2.24x</b>

#### 3.4 Analysis

##### 3.4.1 Memory Access Patterns

**Standard (ijk) - Poor Cache Locality:**

- Innermost loop accesses  $B[k][j]$  with stride N (column-wise)
- Each iteration of k jumps to a different row in B
- For  $512 \times 512$  double matrix: stride =  $512 \times 8$  bytes = 4096 bytes
- This exceeds typical L1 cache line size, causing frequent cache misses
- Matrix B is read N times (poor temporal locality)

**Optimized (ikj) - Excellent Cache Locality:**

- Innermost loop accesses  $B[k][j]$  and  $C[i][j]$  with stride 1 (row-wise)
- Sequential access pattern maximizes spatial locality
- $A[i][k]$  is loaded into register  $r$ , eliminating repeated memory access
- Both  $B$  and  $C$  benefit from cache line prefetching

### 3.4.2 Bandwidth Analysis

The **3.36x improvement** in memory bandwidth ( $8.28 \rightarrow 27.79$  GB/s) indicates:

1. Significantly fewer cache misses in the optimized version
2. Better utilization of cache hierarchy (L1, L2, L3)
3. More efficient memory bus usage

### 3.4.3 Computational Complexity

Both implementations perform the same number of operations:

$$\text{FLOPs} = 2N^3 = 2 \times 512^3 = 268,435,456 \text{ operations}$$

However, the optimized version executes faster due to superior memory access patterns, demonstrating that **memory performance often dominates total execution time**.

## 3.5 Conclusions

1. **Loop Order Matters:** Changing loop order from  $ijk$  to  $ikj$  yields 2.24x speedup without changing the algorithm
2. **Cache-Aware Programming:** Understanding data layout and access patterns is crucial for performance
3. **Register Reuse:** Storing  $A[i][k]$  in a register eliminates redundant memory accesses
4. **Memory is the Bottleneck:** The 3.36x bandwidth improvement demonstrates that memory access, not computation, limits performance in naive implementations

## 4 Exercise 3: Block Matrix Multiplication

### 4.1 Objective

Implement and analyze block (tiled) matrix multiplication to optimize cache usage through data locality.

### 4.2 Methodology

The implementation divides matrices into blocks of size  $B \times B$  and processes them sequentially:

```
1 for (int i0 = 0; i0 < n; i0 += tileSize)
2     for (int j0 = 0; j0 < n; j0 += tileSize)
3         for (int k0 = 0; k0 < n; k0 += tileSize)
4             // Process B x B block
5             for (int i = i0; i < i0 + tileSize && i < n; i++)
6                 for (int k = k0; k < k0 + tileSize && k < n; k++)
7                     for (int j = j0; j < j0 + tileSize && j < n; j++)
8                         C[i*n+j] += A[i*n+k] * B[k*n+j];
```

Tested tile sizes: 16, 32, 64, 128 on a  $512 \times 512$  matrix.

### 4.3 Results

Table 3: Block Matrix Multiplication Performance (N=512)

Tile Size	Time (s)	Bandwidth (GB/s)
16	0.123	0.068
32	0.114	0.074
<b>64</b>	<b>0.101</b>	<b>0.083</b>
128	0.117	0.072

### 4.4 Analysis

#### 4.4.1 Optimal Block Size

Best Performance: Tile Size = 64

- Execution time: 0.101 seconds
- Memory bandwidth: 0.083 GB/s
- 21.9% faster than tile size 16
- 15.8% faster than tile size 128

#### 4.4.2 Cache Hierarchy Considerations

For a typical CPU cache hierarchy:

- L1 Cache: 32-64 KB per core
- L2 Cache: 256-512 KB per core
- L3 Cache: Several MB (shared)

**Memory Requirements per Tile:**

$$\text{Memory} = 3 \times B^2 \times 8 \text{ bytes} = 24B^2 \text{ bytes}$$

Table 4: Tile Memory Footprint

Tile Size	Memory Required
16	6 KB
32	24 KB
64	96 KB
128	384 KB

#### 4.4.3 Why Tile Size 64 is Optimal

##### 1. L1 Cache Fit (16, 32): Too small

- Fits in L1 but creates excessive overhead from outer loop iterations
- More block boundary operations
- Number of tiles:  $(512/16)^3 = 32,768$  vs  $(512/64)^3 = 512$

##### 2. Sweet Spot (64):

- 96 KB fits comfortably in L2 cache (256-512 KB)
- Good balance between cache reuse and loop overhead
- Sufficient parallelism for modern CPUs

##### 3. Too Large (128):

- 384 KB exceeds L2 cache capacity
- Cache thrashing between the three matrices
- Reduced temporal locality

### 4.5 Conclusions

1. **Optimal Block Size = 64:** Provides best performance by fitting in L2 cache
2. **Cache-Aware Blocking Works:** Reduces cache misses through better temporal locality
3. **Trade-offs Exist:** Too small creates overhead; too large causes cache thrashing
4. **Hardware Specific:** Optimal block size depends on cache hierarchy
5. **Bandwidth Metric Caveat:** Low bandwidth can indicate efficient cache reuse

## 5 Exercise 4: Memory Management and Valgrind

### 5.1 Objective

Use Valgrind's Memcheck tool to detect and fix memory leaks in a C program.

### 5.2 Initial Program Analysis

The program allocates two integer arrays:

- **array**: Original array (SIZE = 5 integers = 20 bytes)
- **array\_copy**: Duplicate of the original array (20 bytes)

**Intentional Bug:** The `free_memory()` function was incomplete, and `array_copy` was never freed.

### 5.3 Valgrind Results - Before Fix

```
1 ==15075==  HEAP SUMMARY:
2 ==15075==    in use at exit: 40 bytes in 2 blocks
3 ==15075==    total heap usage: 3 allocs, 1 frees, 1,064 bytes allocated
4 ==15075==
5 ==15075==  20 bytes in 1 blocks are definitely lost in loss record 1 of
6           2
7 ==15075==    at 0x4846828: malloc (vgpreload_memcheck-amd64-linux.so)
8 ==15075==    by 0x109208: allocate_array (memory_debug.c:8)
9 ==15075==    by 0x1093CC: main (memory_debug.c:48)
10 ==15075==
11 ==15075==  20 bytes in 1 blocks are definitely lost in loss record 2 of
12           2
13 ==15075==    at 0x4846828: malloc (vgpreload_memcheck-amd64-linux.so)
14 ==15075==    by 0x109349: duplicate_array (memory_debug.c:34)
15 ==15075==    by 0x109403: main (memory_debug.c:51)
16 ==15075== LEAK SUMMARY:
==15075==  definitely lost: 40 bytes in 2 blocks
```

#### Identified Leaks:

1. 20 bytes from `allocate_array()` call at line 48
2. 20 bytes from `duplicate_array()` call at line 51

### 5.4 Fix Applied

Modified `free_memory()` function:

```
1 void free_memory(int *arr) {
2     free(arr); // Added this line
3 }
```

Modified `main()` function:

```

1 int main() {
2     int *array = allocate_array(SIZE);
3     initialize_array(array, SIZE);
4     print_array(array, SIZE);
5     int *array_copy = duplicate_array(array, SIZE);
6     print_array(array_copy, SIZE);
7     free_memory(array);
8     free_memory(array_copy); // Added this line
9     return 0;
10 }
```

## 5.5 Valgrind Results - After Fix

```

1 ==15094==  HEAP SUMMARY:
2 ==15094==    in use at exit: 0 bytes in 0 blocks
3 ==15094==    total heap usage: 3 allocs, 3 frees, 1,064 bytes allocated
4 ==15094==
5 ==15094== All heap blocks were freed -- no leaks are possible
6 ==15094==
7 ==15094== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from
0)
```

## 5.6 Analysis

### 5.6.1 Memory Leak Types

Valgrind classifies leaks as:

- **Definitely lost:** Memory that is no longer accessible (neither freed nor pointed to)
- **Indirectly lost:** Memory pointed to by definitely lost memory
- **Possibly lost:** Memory that might still be accessible via unusual pointer arithmetic
- **Still reachable:** Memory still pointed to at program exit (not technically a leak)

Our program had **definitely lost** leaks—the most severe type.

### 5.6.2 Valgrind's Detection Mechanism

Valgrind works by:

1. Intercepting all memory allocation/deallocation calls
2. Maintaining a shadow data structure tracking allocation state
3. Recording call stacks for each allocation
4. Analyzing memory state at program exit
5. Reporting unreachable allocated memory with stack traces

## 5.7 Best Practices Demonstrated

1. **Every malloc needs a free:** Each allocation must have corresponding deallocation
2. **Use Valgrind regularly:** Detects leaks that may not cause immediate problems
3. **Fix leaks systematically:** Use stack traces to locate allocation sites
4. **Verify fixes:** Re-run Valgrind to confirm all leaks are resolved

## 5.8 Conclusions

1. **Valgrind is essential:** Detects memory leaks that are invisible at runtime
2. **Memory leaks accumulate:** Even small leaks (40 bytes) are problematic in long-running programs
3. **Prevention is key:** Following RAII patterns and clear ownership prevents leaks
4. **Zero errors achieved:** The fixed program has no memory leaks or errors

## 6 Exercise 5: HPL Benchmark Analysis

### 6.1 Objective

Analyze HPL (High-Performance Linpack) benchmark performance across different matrix sizes (N) and block sizes (NB) on a single core.

### 6.2 Hardware Specifications

**Toubkal Supercomputer Node:**

- Processor: Intel Xeon Platinum 8276L
- Sockets: 2
- Cores per socket: 28
- Base frequency: 2.2 GHz
- Vector instructions: AVX-512 with FMA
- Operations per cycle (per core): 32 DP FLOPs

**Theoretical Peak Performance (Single Core):**

$$P_{\text{core}} = 1 \text{ core} \times 2.2 \times 10^9 \text{ Hz} \times 32 \text{ FLOPs/cycle} = 70.4 \text{ GFLOP/s}$$

### 6.3 Experimental Setup

- MPI processes: 1
- Process grid: P=1, Q=1
- OpenMP threads: 1
- Matrix sizes (N): 1000, 5000, 10000, 20000
- Block sizes (NB): 1, 2, 4, 8, 16, 32, 64, 128, 256
- Total executions: 36

### 6.4 Results

#### 6.4.1 Performance Summary by Matrix Size

Table 5: Best Performance for Each Matrix Size

N	Best NB	Time (s)	GFLOP/s	Efficiency (%)
1000	256	0.55	1.21	1.72%
5000	128	1.81	46.09	65.47%
10000	256	12.73	52.36	74.38%
20000	128	96.83	55.09	78.26%

### 6.4.2 Complete Performance Matrix

Table 6: HPL Performance (GFLOP/s) for All Matrix and Block Sizes

N	NB=1	NB=2	NB=4	NB=8	NB=16	NB=32	NB=64	NB=128	NB=256
1000	0.036	0.067	0.117	0.225	0.402	0.618	0.884	<b>1.168</b>	<b>1.209</b>
5000	0.931	4.753	8.180	15.837	27.705	37.877	44.163	<b>46.092</b>	45.496
10000	1.844	4.505	7.102	14.115	24.827	37.310	45.541	50.559	<b>52.361</b>
20000	1.516	1.968	7.108	4.217	27.381	39.510	49.599	<b>55.087</b>	19.629

Note: Bold values indicate best performance for each matrix size.

### 6.4.3 Detailed Results: N=10000

Table 7: Performance vs Block Size for N=10000

NB	Time (s)	GFLOP/s	Efficiency (%)
1	361.68	1.84	2.62%
2	148.00	4.51	6.40%
4	93.89	7.10	10.09%
8	47.24	14.12	20.05%
16	26.86	24.83	35.27%
32	17.87	37.31	53.00%
64	14.64	45.54	64.69%
128	13.19	50.56	71.82%
<b>256</b>	<b>12.73</b>	<b>52.36</b>	<b>74.38%</b>

## 6.5 Analysis

### 6.5.1 Effect of Matrix Size (N)

#### Performance Scaling:

1. **N=1000:** Very poor efficiency (1.72%)
  - Small problem size  $\Rightarrow$  low computational intensity
  - Overhead dominates: function calls, loop setup, cache initialization
  - Total FLOPs:  $\frac{2}{3}N^3 \approx 0.67$  GFLOPs
2. **N=5000:** Moderate efficiency (65.47%)
  - Problem size allows better cache utilization
  - Computation starts to dominate over overhead
  - Total FLOPs:  $\frac{2}{3} \times 5000^3 \approx 83.33$  GFLOPs
3. **N=10000:** Good efficiency (74.38%)
  - Larger working set keeps processor busy

- Better amortization of initialization costs
- Total FLOPs:  $\frac{2}{3} \times 10000^3 \approx 666.67$  GFLOPs

4. **N=20000:** Best efficiency (78.26%)

- Computation fully dominates
- Approaches realistic scientific computing workloads
- Total FLOPs:  $\frac{2}{3} \times 20000^3 \approx 5333.33$  GFLOPs

**General Trend:** Efficiency increases with problem size, asymptotically approaching but never reaching 100%.

### 6.5.2 Effect of Block Size (NB)

**Performance vs NB (for N=10000):**

- NB=1: 1.84 GFLOP/s (catastrophically poor)
- NB=256: 52.36 GFLOP/s (28.4x improvement)

**Why Small Block Sizes Perform Poorly:**

1. **Excessive function calls:** More panels  $\Rightarrow$  more factorization calls
2. **Poor cache reuse:** Small blocks don't exploit temporal locality
3. **Suboptimal BLAS operations:** BLAS routines optimized for larger blocks
4. **Increased overhead:** Panel pivot search, broadcast, update operations scale with number of panels

**Optimal Block Size Selection:**

For most matrix sizes, NB=128 or NB=256 performs best because:

- Balances cache utilization with computational granularity
- Aligns well with BLAS library optimizations (typically tuned for 64-256)
- Provides sufficient work per panel to amortize overhead
- Matches typical L2/L3 cache sizes

**Exception: N=20000, NB=256**

- Time: 271.73s (worse than NB=128: 96.83s)
- Performance: 19.63 GFLOP/s (vs 55.09 GFLOP/s) — 64.4% performance loss
- Cause: Only 78 panels ( $\lfloor 20000/256 \rfloor$ ) — insufficient parallelism granularity
- NB=128 provides 156 panels, better load balancing for panel operations

**Anomaly: N=20000, NB=8**

- Time: 1264.89s, Performance: 4.22 GFLOP/s

- This is **worse** than NB=4 (7.11 GFLOP/s) and NB=2 (1.97 GFLOP/s)
- Probable cause: Pathological cache behavior or excessive panel operations
- With NB=8, there are 2500 panels — overhead dominates computation
- The data shows that very small blocks (NB < 16) are catastrophically bad for large matrices

### 6.5.3 Efficiency Analysis

**Why Measured Performance  $\ll$  Theoretical Peak (70.4 GFLOP/s)?**

**1. Memory Bandwidth Limitation:**

- LU decomposition has arithmetic intensity  $\approx O(N)$  operations per byte
- At large N, still memory-bound for single core
- Theoretical peak assumes perfect data availability

**2. Algorithmic Constraints:**

- HPL uses right-looking LU factorization with partial pivoting
- Panel factorization is sequential (limits parallelism)
- Pivot search and row swaps introduce dependencies

**3. Cache Hierarchy Effects:**

- Working set eventually exceeds all cache levels
- Main memory latency: 50-100 ns
- Even with prefetching, memory stalls occur

**4. Instruction Mix:**

- Not all instructions are FMA operations
- Address calculations, loads, stores don't contribute to FLOP count
- Branch mispredictions waste cycles

**5. Clock Frequency:**

- 2.2 GHz is base frequency; actual may be lower under sustained load
- Thermal throttling possible on single-core sustained workload

**Realistic Expectations:**

Achieving 70-80% of peak on HPL with a single core is considered excellent. The observed 78.26% efficiency at N=20000 is outstanding and indicates:

- Excellent BLAS library tuning (OpenBLAS)
- Good block size selection
- Effective cache utilization
- Well-optimized panel operations

## 6.6 Performance Formulas

Total Operations for LU Factorization:

$$\text{FLOPs} = \frac{2}{3}N^3 + O(N^2)$$

Efficiency:

$$\eta = \frac{P_{\text{HPL}}}{P_{\text{core}}} = \frac{\text{Measured GFLOP/s}}{70.4 \text{ GFLOP/s}}$$

Example (N=20000, NB=128):

$$\eta = \frac{55.09}{70.4} = 0.7826 = 78.26\%$$

## 6.7 Conclusions

### 1. Matrix Size Impact:

- N=1000: Only 1.72% efficiency — overhead dominates
- N=5000: 65.47% efficiency — computation starts to dominate
- N=10000: 74.38% efficiency — good cache utilization
- N=20000: 78.26% efficiency — excellent performance
- Efficiency grows with problem size but plateaus near 78-80%

### 2. Block Size Optimization:

- NB=1-8: Catastrophically poor (1.5-14 GFLOP/s for large N)
- NB=16-32: Moderate performance (24-40 GFLOP/s)
- NB=64-128: Excellent performance (45-55 GFLOP/s)
- NB=256: Optimal for N=10000; suboptimal for N=20000
- Rule of thumb:  $64 \leq NB \leq 128$  for best results

### 3. Performance Variability:

- Block size can cause 28x performance difference (N=10000: 1.84 vs 52.36 GFLOP/s)
- Poor block size selection is worse than poor matrix size selection
- Anomalies exist: N=20000, NB=8 performs worse than NB=4 (likely cache thrashing)

### 4. Best Configuration: N=20000, NB=128 achieves:

- 55.09 GFLOP/s (78.26% of theoretical peak)
- This is outstanding for single-core HPL
- Demonstrates excellent BLAS library tuning (OpenBLAS)

### 5. Practical Guidelines:

- Always use  $NB \geq 64$  for production workloads
- Test  $NB=128$  and  $NB=256$ ; choose the better performer
- Larger matrices benefit more from optimization
- Single-core efficiency plateaus around 78-80% for HPL