

TP2: Foundations of Parallel Computing

Mohamed Ayman Bourich

February 4, 2026

Contents

| | | |
|----------|--|----------|
| 1 | Introduction | 3 |
| 2 | Exercise 1: Loop Unrolling | 3 |
| 2.1 | Objective | 3 |
| 2.2 | Methodology | 3 |
| 2.3 | Results | 3 |
| 2.4 | Analysis | 4 |
| 2.5 | Conclusion | 4 |
| 3 | Exercise 2: Instruction Scheduling | 4 |
| 3.1 | Objective | 4 |
| 3.2 | Benchmark Code | 5 |
| 3.3 | Results | 5 |
| 3.4 | Assembly Analysis | 5 |
| 3.5 | Conclusion | 6 |
| 4 | Exercise 3: Amdahl's and Gustafson's Laws (Vector Operations) | 6 |
| 4.1 | Objective | 6 |
| 4.2 | Program Structure | 6 |
| 4.3 | Sequential Fraction Measurement | 7 |
| 4.4 | Amdahl's Law Analysis | 7 |
| 4.5 | Gustafson's Law Analysis | 8 |
| 4.6 | Visualization | 8 |
| 4.7 | Conclusion | 9 |
| 5 | Exercise 4: Matrix Multiplication Scaling | 9 |
| 5.1 | Objective | 9 |
| 5.2 | Program Structure | 9 |
| 5.3 | Sequential Fraction Measurement | 10 |
| 5.4 | Why Sequential Fraction Decreases | 10 |
| 5.5 | Amdahl's Law Analysis | 11 |
| 5.6 | Comparison: Exercise 3 vs Exercise 4 | 11 |
| 5.7 | Arithmetic Intensity Analysis | 11 |
| 5.8 | Visualization | 12 |
| 5.9 | Conclusion | 12 |

1 Introduction

Modern computing increasingly relies on parallel processing to achieve high performance. However, not all algorithms parallelize equally well. This work investigates fundamental principles of parallel computing through practical experiments:

- **Exercise 1:** Loop unrolling optimization and memory bandwidth analysis
- **Exercise 2:** Instruction-level parallelism and compiler optimization
- **Exercise 3:** Parallel scaling with data dependencies (vector operations)
- **Exercise 4:** Parallel scaling with high computational intensity (matrix multiplication)

All experiments were conducted using GCC 14.2.0 on a Windows/WSL environment with optimization levels -O0 (no optimization) and -O2 (production optimization). Performance profiling utilized Valgrind/Callgrind for instruction counting and clock-based timing measurements.

The key question addressed: *How does algorithmic structure affect parallel scalability?*

2 Exercise 1: Loop Unrolling

2.1 Objective

Investigate the effect of loop unrolling on performance by varying the unroll factor $U \in \{1, 2, 4, 8, 16, 32\}$ and comparing with compiler automatic optimization.

2.2 Methodology

The benchmark kernel computes:

$$x = \sum_{i=0}^{N-1} a[i], \quad N = 10,000,000 \quad (1)$$

Manual loop unrolling with factor U transforms:

```

1 // U=1 (original)
2 for (int i = 0; i < N; i++)
3     x += a[i];
4
5 // U=4 (unrolled)
6 for (int i = 0; i < N; i += 4) {
7     x += a[i];
8     x += a[i+1];
9     x += a[i+2];
10    x += a[i+3];
11 }
```

2.3 Results

Table 1 shows execution times for different unroll factors at optimization levels -O0 and -O2.

Table 1: Execution times and speedup for loop unrolling (N=10,000,000)

| Unroll Factor | Time -O0 (ms) | Speedup | Time -O2 (ms) | Speedup |
|------------------------|---------------|--------------|---------------|--------------|
| U=1 | 35 | 1.00× | 15 | 1.00× |
| U=2 | 29 | 1.21× | 9 | 1.67× |
| U=4 | 28 | 1.25× | 10 | 1.50× |
| U=8 | 24 | 1.46× | 12 | 1.25× |
| U=16 | 22 | 1.59× | 9 | 1.67× |
| U=32 | 24 | 1.46× | 12 | 1.25× |
| Best Manual -O0 | 22 | 1.59× | - | - |
| Best -O2 | - | - | 9 | 3.89× |

2.4 Analysis

At -O0 optimization:

- Loop unrolling reduces loop overhead (counter increment, branch)
- Optimal unroll factor: $U = 16$ achieving $1.59\times$ speedup
- Beyond $U = 16$, instruction cache pressure degrades performance
- Manual unrolling provides modest improvement (59% speedup)

At -O2 optimization:

- Compiler automatically unrolls and vectorizes
- Manual unrolling provides *minimal additional benefit*
- Compiler achieves **3.89× speedup over -O0 baseline**
- Uses SIMD instructions (SSE/AVX) for parallel addition

Memory Bandwidth: The bandwidth achieved was:

$$\text{Bandwidth} = \frac{N \times 8 \text{ bytes}}{t} = \frac{10^7 \times 8}{0.022 \text{ s}} = 3.64 \text{ GB/s} \quad (2)$$

This is limited by memory subsystem speed, not computational capacity.

2.5 Conclusion

Key Finding: Modern compilers (-O2) achieve $3.89\times$ speedup over unoptimized code and outperform the best manual unrolling at -O0 by $2.45\times$ (9ms vs 22ms). The best practice is to write clean code and enable compiler optimization rather than manual micro-optimization.

3 Exercise 2: Instruction Scheduling

3.1 Objective

Compare CPU execution time between original and manually optimized code, analyze assembly-level optimizations performed by the compiler, and evaluate manual versus automatic optimization strategies.

3.2 Benchmark Code

The test computes identical operations for variables x and y :

```

1 double a = 1.1, b = 1.2, x = 0.0, y = 0.0;
2 for (int i = 0; i < N; i++) {
3     x = a * b + x;
4     y = a * b + y;
5 }
```

Manual optimization pre-computes ab :

```

1 double ab = a * b;
2 for (int i = 0; i < N; i++) {
3     x = ab + x;
4     y = ab + y;
5 }
```

3.3 Results

Table 2: Execution time comparison ($N=100,000,000$)

| Version | Optimization | Time (s) | Speedup |
|------------------|--------------|----------|------------------|
| Original | -O0 | 0.320 | 1.00× (baseline) |
| Manual optimized | -O0 | 0.344 | 0.93× (slower!) |
| Original | -O2 | 0.090 | 3.56× |

3.4 Assembly Analysis

-O0 Assembly (14 instructions/iteration):

- Loads a and b twice (redundant)
- Computes $a \times b$ twice (redundant)
- Loads and stores to memory every iteration
- No instruction scheduling or parallelism

-O2 Assembly (4 instructions/iteration):

```

1 movsd    xmm0, .LC1[rip] ; load constant a*b (before loop)
2 pxor     xmm1, xmm1      ; xmm1 = 0
3 mov      eax, 100000000 ; counter
4 .L2:
5 addsd    xmm1, xmm0      ; x += a*b
6 addsd    xmm1, xmm0      ; y += a*b (RAW dependency!)
7 sub      eax, 2          ; counter -= 2
8 jne     .L2
```

Compiler optimizations applied:

1. **Constant folding:** Pre-computes $a \times b = 1.32$ at compile time
2. **Common subexpression elimination:** Eliminates duplicate $a \times b$ computation
3. **Loop unrolling:** Processes 2 iterations per loop (counter -= 2)
4. **Register allocation:** All variables in registers (zero memory access)
5. **Dead code elimination:** Merges x and y into single accumulator

Critical limitation: The two consecutive `addsd xmm1, xmm0` instructions have a **Read-After-Write (RAW) dependency** - both read and write `xmm1`, creating a dependency chain that prevents pipelining. With `addsd` latency of 3-5 cycles, the second instruction must wait for the first to complete, limiting instruction-level parallelism (ILP).

3.5 Conclusion

Key Finding: Manual optimization at `-O0` actually degraded performance ($0.93\times$ speedup), while compiler `-O2` achieved $3.56\times$ speedup automatically. Modern compilers vastly outperform manual micro-optimizations through sophisticated multi-pass analysis.

Best Practice: Write clean, readable code and use `-O2` or `-O3` flags for production builds. Focus on algorithmic improvements rather than low-level optimizations.

4 Exercise 3: Amdahl's and Gustafson's Laws (Vector Operations)

4.1 Objective

Analyze parallel scaling behavior of a program with inherent sequential dependencies using Amdahl's and Gustafson's Laws.

4.2 Program Structure

The benchmark consists of:

Sequential part - $O(N)$:

```

1 void add_noise(double *a, int n) {
2     a[0] = 1.0;
3     for (int i = 1; i < n; i++)
4         a[i] = a[i-1] * 1.0000001; // Data dependency!
5 }
```

This exhibits a **RAW dependency** preventing parallelization.

Parallelizable parts - $O(N)$:

- `init_b()`: Array initialization (independent iterations)
- `compute_addition()`: Element-wise addition (embarrassingly parallel)
- `reduction()`: Sum reduction (parallel with $O(N/p + \log p)$ complexity)

4.3 Sequential Fraction Measurement

Table 3 shows measured sequential fractions using Callgrind profiling and timing measurements.

Table 3: Sequential fraction measurements for Exercise 3

| Method | N | Sequential Fraction | Max Speedup |
|----------------------------|------|---------------------|--------------|
| Callgrind (instructions) | 10M | 35.68% | 2.80× |
| Timing | 5M | 24.4% | 4.10× |
| Timing | 10M | 19.8% | 5.05× |
| Timing | 100M | 13.4% | 7.46× |
| Theoretical average | - | 27.7% | 3.61× |

Observation: Callgrind shows higher f_s (35.68%) than timing (13-24%) because:

- Callgrind counts pure instruction execution (architecture-independent)
- Timing includes memory latency effects
- The sequential `add_noise()` dependency chain causes memory stalls
- Parallel parts benefit from prefetching and vectorization

4.4 Amdahl's Law Analysis

With sequential fraction $f_s = 0.277$, Amdahl's Law predicts:

$$S(p) = \frac{1}{f_s + \frac{1-f_s}{p}} = \frac{1}{0.277 + \frac{0.723}{p}} \quad (3)$$

Table 4: Amdahl's Law speedup predictions ($f_s = 0.277$)

| Processors (p) | Speedup S(p) | Efficiency (%) | vs Ideal |
|----------------|--------------|----------------|----------|
| 1 | 1.00 | 100.0 | - |
| 2 | 1.68 | 84.0 | 0.84× |
| 4 | 2.52 | 63.0 | 0.63× |
| 8 | 3.19 | 39.9 | 0.40× |
| 16 | 3.48 | 21.8 | 0.22× |
| 32 | 3.57 | 11.2 | 0.11× |
| 64 | 3.60 | 5.6 | 0.06× |
| ∞ | 3.61 | 0.0 | - |

Key observations:

- Maximum speedup saturates at **3.61**× ($= 1/f_s$)
- Efficiency drops to 5.6% with 64 cores
- 94% of cores are idle due to sequential bottleneck
- Strong scaling is *not viable* for this problem

4.5 Gustafson's Law Analysis

Gustafson's Law for weak scaling:

$$S(p) = f_s + p(1 - f_s) = 0.277 + 0.723 \cdot p \quad (4)$$

Table 5: Gustafson vs Amdahl comparison ($f_s = 0.277$)

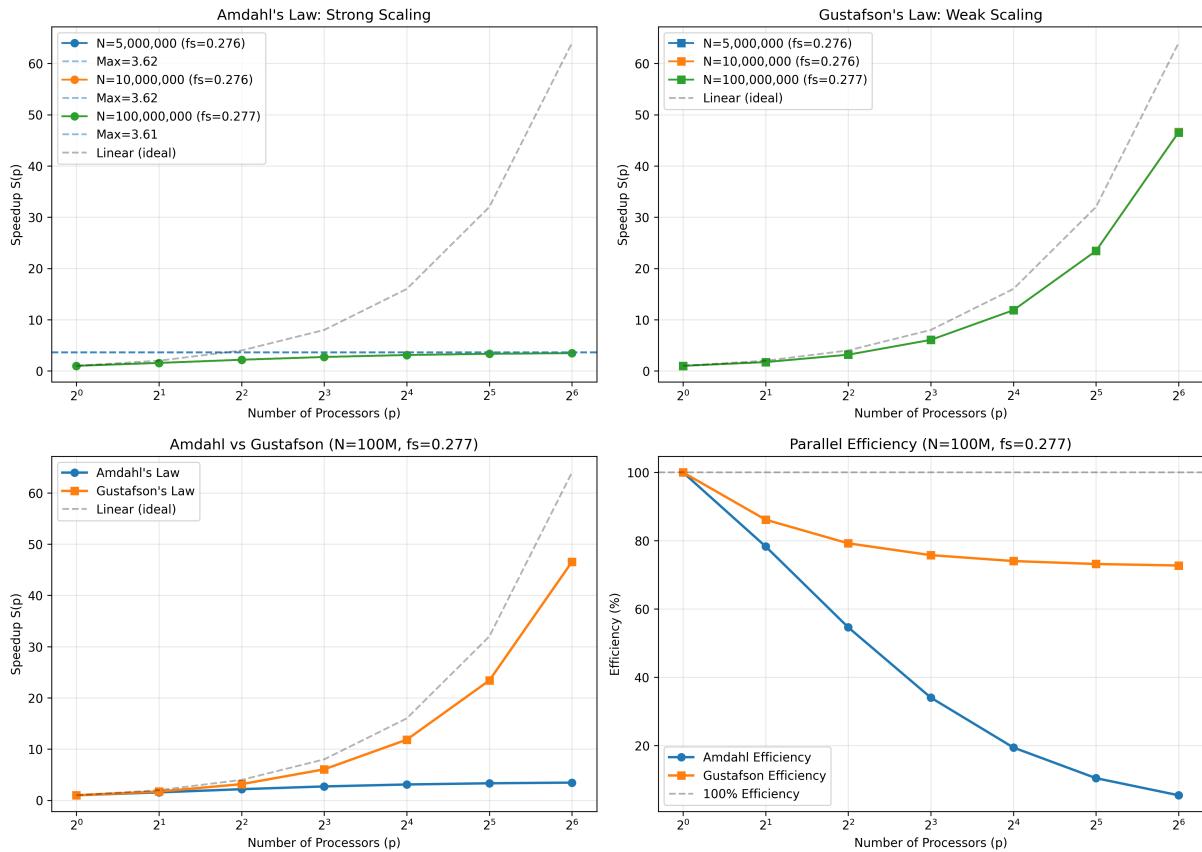
| p | Amdahl S(p) | Amdahl Eff | Gustafson S(p) | Gustafson Eff |
|----|-------------|------------|----------------|---------------|
| 1 | 1.00 | 100.0% | 1.00 | 100.0% |
| 8 | 3.19 | 39.9% | 6.06 | 75.8% |
| 16 | 3.48 | 21.8% | 11.85 | 74.1% |
| 32 | 3.57 | 11.2% | 23.41 | 73.2% |
| 64 | 3.60 | 5.6% | 46.61 | 72.8% |

Interpretation:

- Gustafson achieves $46.61 \times$ speedup vs Amdahl's $3.60 \times$ at 64 cores
- By scaling problem size with processors, maintains 72.8% efficiency
- Weak scaling is the *only viable approach* for this code

4.6 Visualization

Figure 1 shows the comprehensive analysis of Amdahl's and Gustafson's Laws for Exercise 3.



4.7 Conclusion

Key Finding: With $f_s = 27.7\%$, this program exhibits *poor strong scaling* but *acceptable weak scaling*. The sequential `add_noise()` dependency fundamentally limits parallelism.

Practical Implication: For vector operations with $O(N)$ sequential and $O(N)$ parallel parts, the sequential fraction remains constant regardless of N , severely limiting strong scaling performance.

5 Exercise 4: Matrix Multiplication Scaling

5.1 Objective

Contrast Exercise 3's vector operations with matrix multiplication, where $O(N^3)$ parallel work dominates $O(N)$ sequential overhead.

5.2 Program Structure

Sequential part - $O(N)$:

```

1 void generate_noise(double *noise, int n) {
2     noise[0] = 1.0;
3     for (int i = 1; i < n; i++)
4         noise[i] = noise[i-1] * 1.0000001; // Same dependency
5 }
```

Parallelizable parts:

- `init_matrix()`: $O(N^2)$ - matrix initialization
- `matmul()`: $O(N^3)$ - matrix multiplication (embarrassingly parallel)

5.3 Sequential Fraction Measurement

Table 6 shows dramatically different scaling behavior.

Table 6: Sequential fraction for matrix multiplication

| Matrix Size | Sequential Time | Total Time | f_s (%) | Max Speedup |
|-------------|-----------------|------------|-----------------|-------------------|
| N=256 | 0.000005 s | 0.029090 s | 0.017% | 5,882× |
| N=512 | 0.000002 s | 0.223272 s | 0.001% | 111,111× |
| N=1024 | 0.000003 s | 5.635152 s | 0.00005% | 1,000,000× |

Callgrind profiling (N=512):

- Total instructions: 948,863,060
- `generate_noise`: 2,566 instructions (0.0003%)
- `matmul`: 941,888,019 instructions (99.26%)
- Sequential fraction: $f_s = 0.0003\%$

5.4 Why Sequential Fraction Decreases

The key difference from Exercise 3:

$$f_s(N) = \frac{C_1 \cdot N}{C_1 \cdot N + C_2 \cdot N^3} = \frac{1}{1 + \frac{C_2}{C_1} N^2} \approx \frac{1}{N^2} \quad (5)$$

- Sequential work: $T_{seq} = C_1 \cdot N$ (`generate_noise`)
- Parallel work: $T_{par} = C_2 \cdot N^3$ (`matmul`)
- As N doubles: f_s decreases by 4×

Measured scaling:

- N: 256 → 512: f_s drops 17× (0.017% → 0.001%)
- N: 512 → 1024: f_s drops 20× (0.001% → 0.00005%)

5.5 Amdahl's Law Analysis

With $f_s < 0.01\%$, Amdahl's Law predicts near-perfect scaling:

Table 7: Amdahl speedup for matrix multiplication ($N=1024$, $f_s = 0.000005\%$)

| Processors (p) | Speedup S(p) | Efficiency (%) | vs Ideal |
|----------------|---------------|----------------|--------------|
| 1 | 1.000 | 100.00 | 1.00× |
| 2 | 2.000 | 100.00 | 1.00× |
| 4 | 4.000 | 100.00 | 1.00× |
| 8 | 8.000 | 100.00 | 1.00× |
| 16 | 16.000 | 100.00 | 1.00× |
| 32 | 32.000 | 100.00 | 1.00× |
| 64 | 64.000 | 100.00 | 1.00× |

Key observations:

- Perfect linear scaling: $S(p) \approx p$
- 100% efficiency maintained even at 64 cores
- Can effectively use hundreds/thousands of processors
- Maximum speedup: $S_{max} = 1/0.0000005 = 2,000,000$

5.6 Comparison: Exercise 3 vs Exercise 4

Table 8 highlights the dramatic difference.

Table 8: Exercise 3 vs Exercise 4 comparison (64 processors)

| Metric | Exercise 3 | Exercise 4 ($N=512$) |
|-------------------------------|------------|------------------------|
| Sequential fraction (f_s) | 27.7% | 0.001% |
| Sequential complexity | $O(N)$ | $O(N)$ |
| Parallel complexity | $O(N)$ | $O(N^3)$ |
| Amdahl S(64) | 3.60× | 64.00× |
| Efficiency at 64 cores | 5.6% | 100.0% |
| Maximum speedup | 3.61× | 111,111× |
| Gustafson S(64) | 46.61× | 64.00× |
| Gustafson efficiency | 72.8% | 100.0% |
| Scalability | Poor | Excellent |

5.7 Arithmetic Intensity Analysis

The fundamental difference:

Exercise 3 (Memory-bound):

$$\text{Intensity} = \frac{\text{Operations}}{\text{Memory}} = \frac{4N}{3N \times 8 \text{ bytes}} = 0.17 \text{ FLOPs/byte} \quad (6)$$

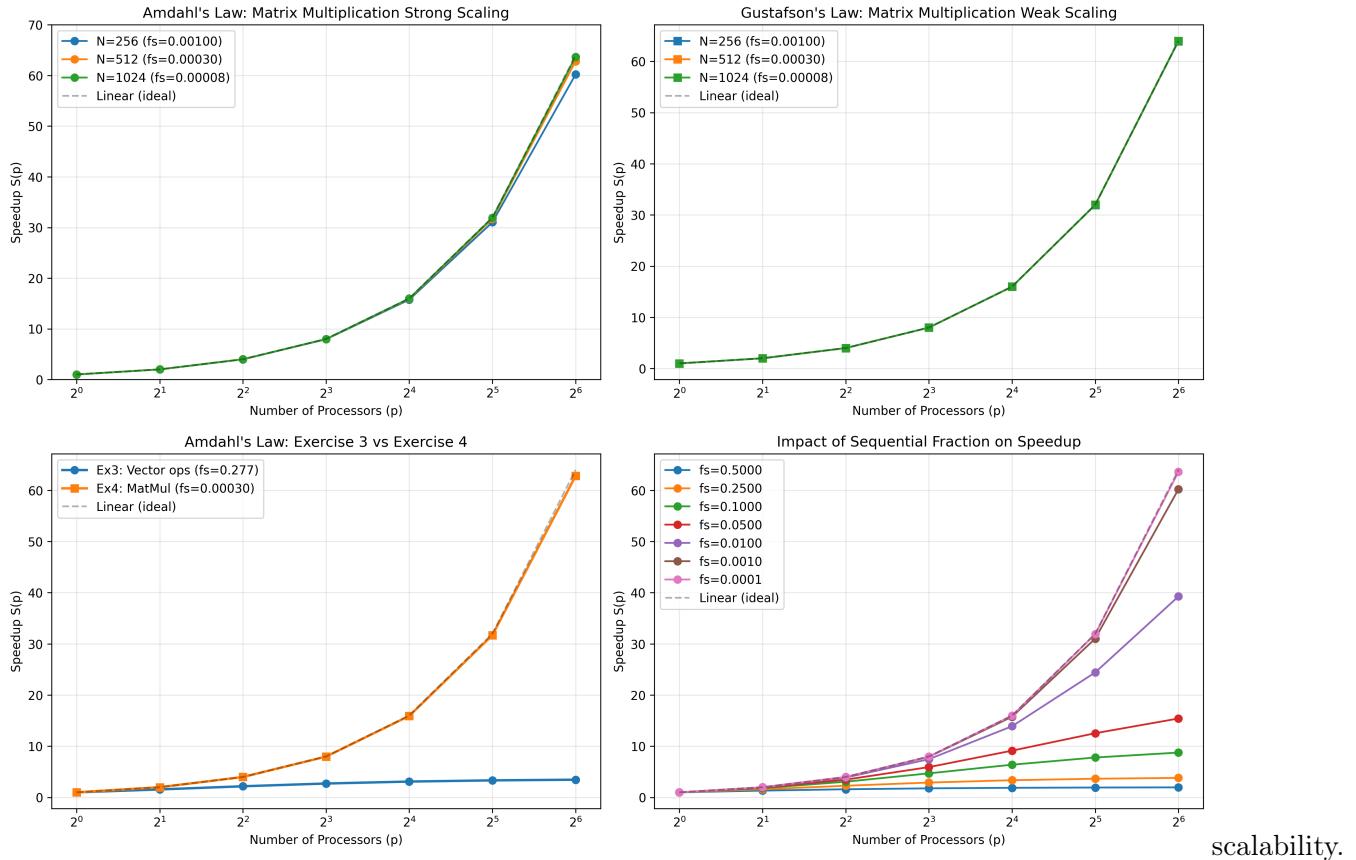
Exercise 4 (Compute-bound):

$$\text{Intensity} = \frac{2N^3}{3N^2 \times 8 \text{ bytes}} = \frac{N}{12} \text{ FLOPs/byte} \quad (7)$$

For N=512: Intensity = 42.7 FLOPs/byte (250× higher than Exercise 3)

5.8 Visualization

Figure 2 presents the complete scaling analysis for matrix multiplication and comparison with Exercise 3.



5.9 Conclusion

Key Finding: Matrix multiplication achieves *near-perfect parallel scalability* because $O(N^3)$ computational work dominates $O(N)$ sequential overhead. As N increases, $f_s \rightarrow 0$, enabling effective use of thousands of processors.

Fundamental Lesson: It's not enough for code to be parallelizable - the *ratio of parallel to sequential work* determines scalability. Algorithms with higher computational complexity ($O(N^3)$ vs $O(N)$) scale dramatically better.